

Model Optimization and Tuning Phase Template

Date	July 2024
Team ID	740103
Project Title	The Language Of Youtube: A Text Classification Approach To Video Descriptions
Maximum Marks	10 Marks

Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase involves refining neural network models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

Hyperparameter Tuning Documentation (8 Marks):

Model	Tuned Hyperparameters
Linear Regression	-
Random Forest Regressor	

<p>Logistic Regression</p>	<p>#importing the library for grid search from sklearn.model selection import GridSearchCV</p> <p>The 'lr_param_grid' specifies different values for regularization strength (C), solvers (solver), and penalty types (penalty). GridSearchCV (lr_cv) is employed with 5-fold cross-validation (cv=5), evaluating model performance based on accuracy (scoring="accuracy"). The process uses all available CPU cores (n_jobs=-1) for parallel processing and provides verbose output (verbose=True) to track progress.</p> <pre>LOGISTIC REGRESSION HYPER PARAMETER TUNNING</pre> <pre>[54] #finding the grid search cv for logistic regression lr=LogisticRegression(n_jobs=-1,random_state=0) lr_param_grid={ 'C':[0.1,0.5,1,5,10], 'solver':['liblinear','saga'], 'penalty':['l1','l2'] } lr_cv=GridSearchCV(lr,lr_param_grid,cv=5,scoring="accuracy",n_jobs=-1,verbose=T lr_cv.fit(x_train,y_train)</pre> <p>Fitting 5 folds for each of 20 candidates, totalling 100 fits /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:1211: warnings.warn(<pre> > GridSearchCV > estimator: LogisticRegression > LogisticRegression </pre> </p>
<p>Decision Tree Regressor</p> <p>SVM</p>	<p>-</p> <p>-</p> <p>-</p>

<p>Random Forest</p>	<p>The parameter <code>grid</code> (<code>rfc_param_grid</code>) for <code>hyperparameter</code> tuning. It specifies different values for the number of trees (<code>n_estimators</code>), splitting criterion (<code>criterion</code>), maximum depth of trees (<code>max_depth</code>), and maximum number of features considered for splitting (<code>max_features</code>). <code>GridSearchCV</code> (<code>rfc_cv</code>) is employed with 3-fold cross-validation (<code>cv=3</code>), evaluating model performance based on accuracy (<code>scoring="accuracy"</code>).</p> <pre> RF= RandomForestClassifier(n_estimators=16,max_depth=130) RF.fit(x_train,y_train) y_pred =RF.predict(x_test) print("Accuracy on test set: %.3f%%"%(accuracy_score(y_test, y_pred)*100)) print("Precision on test set: %.3f"%(precision_score(y_test, y_pred,average='macro')))) print("Recall on test set: %.3f"%(recall_score(y_test, y_pred,average='macro')))) print("F1-Score on test set: %.3f"%(f1_score(y_test, y_pred,average='macro')))) print("-"*20, "confusion matrix", "-"*20) plt.figure(figsize=(12,8)) df_cm = pd.DataFrame(confusion_matrix(y_test, y_pred), range(6),range(6)) sns.set(font_scale=1.4)#for label size labels = ['Art&Music','Food','History','Sci&Tech','Manu','TravelBlog'] sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g',xticklabels=labels, yticklabels=labels) plt.xlabel('Predicted Class') plt.ylabel('Original Class') plt.show() plotPrecisionRecall(y_test,y_pred) </pre>
<p>Decision Tree</p>	<p>The parameters (<code>params</code>) define a grid for hyperparameter tuning of the Decision Tree Classifier (<code>DecisionTreeClassifier</code>), including <code>max_depth</code>, <code>min_samples_leaf</code>, and criterion ('gini' or 'entropy'). <code>GridSearchCV</code> (<code>dec_cv</code>) is used with 5-fold cross-validation (<code>cv=5</code>), evaluating model performance based on accuracy (<code>scoring="accuracy"</code>)</p> <pre> DECISION TREE CLASSIFIER-HYPER PARAMETER TUNNING [68] #finding grid search cv for decision tree classifier dec=DecisionTreeClassifier(random_state=42) params={ 'max_depth': [2, 3, 5, 10, 20], 'min_samples_leaf': [5, 10, 20, 50, 100], 'criterion': ['gini', 'entropy'] } dec_cv=GridSearchCV(dec,param_grid=params,cv=5,n_jobs=-1,scoring="accuracy",verbose=3) dec_cv.fit(x_train,y_train) </pre> <p>Fitting 5 folds for each of 50 candidates, totalling 250 fits</p> <pre> GridSearchCV estimator: DecisionTreeClassifier DecisionTreeClassifier </pre>

SDGClassifier	<p>SGDClassifier is a popular linear classifier in machine learning that uses stochastic gradient descent (SGD) to optimize the loss function. It is well-suited for handling large-scale and sparse datasets, and can be used for both binary and multiclass classification problems. SGDClassifier allows for tuning of hyperparameters such as the learning rate, regularization, and loss function, making it a flexible and powerful tool for classification tasks.</p> <pre> clf = SGDClassifier(loss = 'hinge', alpha = 0.01, class_weight='balanced', learning_rate='optimal',eta0=0.001, n_jobs = -1) clf.fit(x_tr_uni,y_train) y_pred = clf.predict(x_test_uni) print("Accuracy on test set: %0.3f"%(accuracy_score(y_test, y_pred)*100)) print("Precision on test set: %0.3f"%(precision_score(y_test, y_pred,average='macro')))) print("Recall on test set: %0.3f"%(recall_score(y_test, y_pred,average='macro')))) print("F1-Score on test set: %0.3f"%(f1_score(y_test, y_pred,average='macro')))) print("-"*20, "confusion matrix", "-"*20) plt.figure(figsize=(12,8)) matrix=confusion_matrix(y_test, y_pred) df_cm = pd.DataFrame(matrix) sns.set(font_scale=1.4)#for label size labels = ['Art&Music','Food','History','Sci&Tech','Manu','TravelBlog'] sns.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt='g',xticklabels=labels, yticklabels=labels) plt.xlabel('Predicted Class') plt.ylabel('Original Class') plt.show() plotPrecisionRecall(y_test,y_pred) </pre>
---------------	---

Final Model	Reasoning
--------------------	------------------

Final Model Selection Justification (2 Marks):

Linear
SVM
model

The results show that all models have achieved good accuracy, with the Random Forest Classifier using BOW achieving the highest accuracy of 86.478%. However, the SGDClassifier using TF-IDF achieved the highest precision, recall, and F1-score of 0.917. This suggests that the SGDClassifier using TF-IDF is better at correctly identifying positive instances (precision), identifying all positive instances (recall), and balancing both metrics (F1-score) than the other models

Support Vector Machine

Model	hyper parameter	F1-score cv	F1-score test	Precision test	Recall test	Accuracy Test
unigram	0.01	0.927	0.907	0.908	0.907	90.836%
TF-IDF	0.0001	0.922	0.920	0.920	0.920	92.138%

Random Forest

Model	n_estimators	Max Depth	F1-score cv	F1-score test	Precision test	Recall test	Accuracy Test
unigram	16	130	0.868	0.858	0.860	0.858	86.119%
TF-IDF	20	190	0.872	0.849	0.854	0.848	85.085%